

Project Non-ML 5

Final Project for the CM course 2020/21

Matteo De Francesco

Contents

1	Introduction	1
2	Deflected subgradient Analysis	1
2.1	ADAGRAD introduction	1
2.2	Algorithm characteristics	2
2.3	Proximal term discussion	2
2.4	Convergence Analysis	3
2.5	General application to our case	6
3	Implementation	8
4	Code description	12
	References	13
A	Appendix A: Updates derivations	14
A.1	Differentiating proximal term	14
A.2	Derivation of primal-dual update	14
A.3	Derivation of composite-mirror update	14
B	Appendix B: subgradient computation	15

1 Introduction

In this report we will analyze the convex quadratic problem

$$\min \left\{ x^T Q x + q x : \sum_{i \in I^k} x_i = 1, k \in K, x \geq 0 \right\} \quad (P)$$

with the following constraints: $x \in \mathbb{R}^n$, the index sets I^k form a partition of the set $\{1, \dots, n\}$ (i.e. $\cup_{k \in K} I^k = \{1, \dots, n\}$, and $I^h \cap I^k = \emptyset$ for all h and k), and Q is positive semidefinite. The aim of this project is to exploit the Lagrangian dual problem and solve it via one of the *deflected subgradient* methods.

We can identify the inequality and equality constraints and rewrite them in the typical form

$$\begin{aligned} G(x) \rightarrow g_i(x) \leq 0 &\implies -x_i \leq 0 \quad \forall i \in \{1, \dots, n\} \\ H(x) \rightarrow h_j(x) = 0 &\implies \sum_{i \in I^k} x_i - 1 = 0 \quad \forall k \in K \end{aligned}$$

Hence we can rewrite the problem in the following form:

$$\begin{cases} \min & x^T Q x + q x \\ & -x_i \leq 0 \quad \forall i \in \{1, \dots, n\} \\ & \sum_{i \in I^k} x_i - 1 = 0 \quad \forall k \in K \end{cases}$$

In order to have less dual variables, we can rewrite the above problem as a lagrangian function of only the inequality constraints:

$$\mathcal{L}(x; \lambda) = x^T Q x + q^T x + \langle \lambda, x \rangle$$

Now, we can easily construct the lagrangian dual function considering the equality constraints

$$\psi(\lambda) = \min_{x \in \mathcal{Y}} \{ \mathcal{L}(x; \lambda), \lambda \geq 0 \}$$

where $\mathcal{Y} = \{ \sum_{i \in I^k} x_i = 1 \quad \forall k \in K \}$. Hence the following optimization problem

$$\begin{cases} \max_{\lambda} & \psi(\lambda) \\ \text{subject to} & \lambda \geq 0 \end{cases} \quad (D)$$

We are assuming that optimizing over the set \mathcal{Y} can be done very easily.

In the next section, we will briefly recall the properties of the **ADAGRAD** family of algorithms[1].

2 Deflected subgradient Analysis

2.1 ADAGRAD introduction

The projection rule of a point y onto the constraint set \mathcal{X} according to a positive semidefinite matrix A amounts to:

$$\prod_{\mathcal{X}}^A(y) = \arg \min_{\lambda \in \mathcal{X}} \|\lambda - y\|_A = \arg \min_{\lambda \in \mathcal{X}} \langle \lambda - y, A(\lambda - y) \rangle$$

and it's aimed to minimize the Mahalanobis norm.

ADAGRAD applies the following projection rule:

$$\lambda_{t+1} = \prod_{\mathcal{X}}^{diag(G_t)^{1/2}} (\lambda_t - \eta diag(G_t)^{-1/2} g_t) \quad (1)$$

where the matrix A coincides with the diagonal square root of the full outer product of the subgradients $G_t = \sum_{\tau=1}^t g_{\tau} g_{\tau}^T$. As we stated in the introduction, we should exploit **ADAGRAD**

on the dual problem, whose constraint set consists in the simple inequality constraint $\lambda \geq 0$, which is addressable by ADAGRAD, given the fact that the algorithm can be applied to any convex set $\mathcal{X} \subseteq \mathbb{R}^n$, respected by our problem (we are in \mathbb{R}_+^n).

In addition, as the general projection rule tell us, the new update λ_{t+1} is projected over the constraint set \mathcal{X} according to the matrix of the gradients.

We will recall in the next section some algorithmic properties of the algorithm convergence for diagonal matrices, and we will derive the update rule followed by the projection over the constraint set $\mathcal{X} = \{\lambda \geq 0\}$.

2.2 Algorithm characteristics

First of all, we reiterate that ADAGRAD is suitable for the dual problem, given its application in a convex setting. The goal is to attain a small regret bound:

$$R_\phi(T) \triangleq \sum_{t=1}^T \psi(\lambda_t) - \psi(\lambda^*) \quad (2)$$

between the actual value of the dual function ψ with the corresponding iterate λ at step t and the value of ψ with the optimal solution λ^* .

The point of ADAGRAD is that not all features are equal, hence they must be treated differently. That's the purpose of using an adaption to the geometry of space, so do not use anymore a standard gradient descent but conditioning the different values based on a positive semidefinite matrix A (the G_t in this case).

Two algorithm versions are analyzed in [1], where we omitted the regularization term due to our problem setting. The first update, referred to as *primal-dual subgradient method*, is

$$\lambda_{t+1} = \arg \min_{\lambda \in \mathcal{X}} \{ \eta \langle \bar{g}_t, \lambda \rangle + \frac{1}{t} \Psi_t(\lambda) \} \quad (3)$$

coming from [2], where $\bar{g}_t = \frac{1}{t} \sum_{\tau=1}^t g_\tau$ is the average gradient, η is a fixed stepsize and Ψ_t is the *proximal term*.

The second update instead

$$\lambda_{t+1} = \arg \min_{\lambda \in \mathcal{X}} \{ \eta \langle g_t, \lambda \rangle + B_{\Psi_t}(\lambda, \lambda_t) \} \quad (4)$$

where B refers to the Bregman divergence. This second update comes from [3].

Finally, also the previous mentioned projection rule can be used as an update:

$$\lambda_{t+1} = P_{\mathcal{X}} \{ \lambda_t + \eta \text{diag}(G_t)^{-1/2} g_t \} \quad (5)$$

where P denotes the projection operation. The proximal term Ψ_t is the key point of the algorithm. Instead of using a fixed proximal functions, both the updates use squared Mahalanobis norms as their proximal functions, setting then $\Psi_t(\lambda) = \langle \lambda, H_t \lambda \rangle$ for a symmetric matrix $H_t \succeq 0$. In particular, the diagonal case which we will recall here make use of:

$$H_t = \delta I + \text{diag}(G_t)^{1/2}$$

for some small fixed $\delta \geq 0$.

The usage of a strongly convex proximal function is also remarked in [4] in appendix A.

2.3 Proximal term discussion

In order to attain a lower regret bound and to adapt to the geometry of the space, the objective of the authors is to not use anymore a standard proximal functions but a modified version of it.

The proximal function act as a regularization term, typically the Euclidian projection over a convex set. Given the indicator function of a convex set C :

$$I_C(\lambda) = \begin{cases} 0 & \lambda \in C \\ +\infty & \text{otherwise} \end{cases}$$

the standard proximal mapping of I_C is the Euclidean projection on C :

$$\text{prox}_{I_C}(\lambda) = \arg \min_{u \in C} \|u - \lambda\|_2^2 = P_C(\lambda)$$

Instead in this case the authors noticed that some local region of the function to be optimized need more "attention than others". What they come up with then is the modification (also remarked in the **ADAGRAD** introduction) using a different proximal function where

$$\text{prox}_{I_C}(\lambda) = \arg \min_{u \in C} \|u - \lambda\|_A^2 = P_C(\lambda)$$

the Euclidean projection is not computed anymore according to a 2-norm, but according to a matrix A which in the case of **ADAGRAD** coincide with the matrix of the subgradients. A summary of the properties and aspects of some proximal functions can be found in [5]. Below is reported the analysis of the proximal function regarding [1].

Examining the regret bounds for the updates in [6] and [2], it is quite obvious that they depends on dual norms of the derivative of the function to be optimized, and in turn this last depend on the choice of Ψ . The objective of [1] is to properly modify the value of Ψ during the run of the algorithm in order to lower the contribution of the norms, and so lower the regret bound. This is achieved by keeping second order information about the sequence of iterates λ_t and allow Ψ to vary on each round of the algorithm. To achieve this, we must assume that Ψ_t is monotonically non-decreasing, 1-strongly convex with respect to a time-dependent semi-norm $\|\cdot\|_{\Psi_t}$. Formally, given two generic points x and y , Ψ is 1-strongly convex with respect to $\|\cdot\|_{\Psi}$ if

$$\Psi(y) \geq \Psi(x) + \langle \nabla \Psi(x), y - x \rangle + \frac{1}{2} \|x - y\|_{\Psi}^2$$

As a consequence, strong convexity is guaranteed if and only of $B_{\Psi_t}(x, y) \geq \frac{1}{2} \|x - y\|_{\Psi_t}^2$. The following bound holds then on proximal term, respectively for (3) and (4). Proofs can be found in Appendix F of [1].

Proposition 2.1. *Let the sequence $\{\lambda_t\}$ be defined by the update (3). For any $\lambda^* \in \mathcal{X}$*

$$\sum_{t=1}^T \psi(\lambda_t) - \psi(\lambda^*) \leq \frac{1}{\eta} \Psi_T(\lambda^*) + \frac{\eta}{2} \sum_{t=1}^T \|\psi'(\lambda_t)\|_{\Psi_{t-1}^*}^2 \quad (6)$$

Proposition 2.2. *Let the sequence $\{\lambda_t\}$ be defined by the update (4). For any $\lambda^* \in \mathcal{X}$*

$$\begin{aligned} \sum_{t=1}^T \psi(\lambda_t) - \psi(\lambda^*) &\leq \frac{1}{\eta} B_{\Psi_1}(\lambda^*, \lambda_1) + \frac{1}{\eta} \sum_{t=1}^{T-1} [B_{\Psi_{t+1}}(\lambda^*, \lambda_{t+1}) - B_{\Psi_t}(\lambda^*, \lambda_{t+1})] \\ &\quad + \frac{\eta}{2} \sum_{t=1}^T \|\psi'(\lambda_t)\|_{\Psi_t^*}^2 \end{aligned}$$

The proximal term and the bregman divergence present in (3) and (4) will be computed according to the subgradient matrix, in order to iteratively modify the update and thus adapting to the geometry of the space.

2.4 Convergence Analysis

First of all, let us state that the convergence rate of **ADAGRAD** is the same of the Stochastic Gradient Descent (SGD), hence $O(1/\sqrt{T})$ but with a lower constant due to the use of G matrix.

The algorithm 1 reported in [1] will be applied to our problem and is reported here:

Algorithm 1 ADAGRAD for diagonal matrices

function ADAGRAD(η, δ)

 $x_1 \leftarrow 0$
 $\lambda_1 \leftarrow 0$
 $g_{1:0} \leftarrow []$
for $t \leftarrow 1$ to T **do**
 $Loss = f_t(x_t)$
 $g_t \leftarrow \partial\psi(\lambda_{t-1})$ of ψ at λ_{t-1}
 \triangleright Compute subgradient at λ_{t-1}
 $g_{1:t} \leftarrow [g_{1:t-1} \ g_t]$
 \triangleright Store subgradient

 $s_{t,i} \leftarrow \|g_{1:t,i}\|_2$
 \triangleright Compute the optimal $s_{t,i}$
 $H_t \leftarrow \delta I + \text{diag}(s_t)$
 $\Psi_t(\lambda) \leftarrow \frac{1}{2} \langle \lambda, H_t \lambda \rangle$

Either compute (3) or (4)

end for
end function

The general convergence result of this algorithm for both the updates is reported in theorem 5 of the original paper. We will now report the theoretical analysis behind procedure 1 and lastly modify the algorithm in order to match our problem settings.

Theorem 2.1. *Let the sequence $\{\lambda_t\}$ be defined by algorithm 1. For λ_t generated using the update (3) with $\delta \geq \max_t \|g_t\|_\infty$, for any $\lambda^* \in \mathcal{X}$*

$$R_\phi(T) \leq \underbrace{\frac{\delta}{\eta} \|\lambda^*\|_2^2 + \frac{1}{\eta} \|\lambda^*\|_\infty^2 \sum_{i=1}^d \|g_{1:T,i}\|_2}_{(1)} + \eta \underbrace{\sum_{i=1}^d \|g_{1:T,i}\|_2}_{(2)} \quad (7)$$

For λ_t generated using the update (4), for any $\lambda^* \in \mathcal{X}$

$$R_\phi(T) \leq \underbrace{\frac{1}{2\eta} \max_{t \leq T} \|\lambda^* - \lambda_t\|_\infty^2 \sum_{i=1}^d \|g_{1:T,i}\|_2}_{(3)} + \eta \underbrace{\sum_{i=1}^d \|g_{1:T,i}\|_2}_{(2)} \quad (8)$$

We will now delve into analyzing the different (\odot) parts.

First of all, we should recall some intuition in order to understand better the regret bound provided above.

Focusin on the algorithm 1, the chosen value of $s_{t,i}$ explain us why the particular choice of the proximal function $\Psi_t(x) = \frac{1}{2} \langle x, H_t x \rangle$ is so important. $s_{t,i}$ comes from the solution of the problem

$$\min_s \sum_{t=1}^T \sum_{i=1}^d \frac{g_{t,i}^2}{s_i} \text{ s.t. } s \succeq 0, \langle 1, s \rangle \leq c$$

solved by optimizing the Lagrangian

$$\mathcal{L}(s, \lambda, \theta) = \sum_{i=1}^d \frac{\|g_{1:T,i}\|_2^2}{s_i} - \langle \lambda, s \rangle + \theta(\langle 1, s \rangle - c)$$

Taking the partial derivatives to find the infimum of \mathcal{L} , and using the complementary slackness condition on $\lambda_i s_i$ imply that $\lambda_i = 0$. As a consequence, we obtain $s_i = c \|g_{1:T,i}\|_2 / \sum_{j=1}^d \|g_{1:T,j}\|_2$. Plugging this into the previous objective function, we get

$$\inf_s \left\{ \sum_{t=1}^T \sum_{i=1}^d \frac{g_{t,i}^2}{s_i} : s \succeq 0, \langle 1, s \rangle \leq c \right\} = \frac{1}{c} \left(\sum_{i=1}^d \|g_{1:T,i}\|_2 \right)^2$$

Now it is natural to suspect that for s achieving the infimum in this latter equation, using a proximal function similar to $\Psi(\lambda) = \langle \lambda, \text{diag}(s) \lambda \rangle$ with associated squared dual norm

$\|\lambda\|_{\Psi^*}^2 = \langle \lambda, \text{diag}(s)^{-1} \lambda \rangle$ we should lower the gradient terms both in (7) and (8). The upper bound on the gradient term for both the updates is taken from the LEMMA 4 of [1], stating:

Lemma 2.1. *Let $g_t = \psi'(\lambda_t)$ and $g_{1:t}$ and s_t be defined as in algorithm 1. Then*

$$\sum_{t=1}^T \langle g_t, \text{diag}(s_t)^{-1} g_t \rangle \leq 2 \sum_{i=1}^d \|g_{1:T,i}\|_2$$

To obtain a bound, we need to consider the terms consisting of the dual-norm of the subgradient in the bounds (7) and (8), which is $\|\psi'(\lambda_t)\|_{\Psi_t^*}^2$. When we choose $\Psi_t(\lambda) = \langle \lambda, (\delta I + \text{diag}(s_t)) \lambda \rangle$, the associated dual norm is

$$\|g\|_{\Psi_t^*}^2 = \langle g, (\delta I + \text{diag}(s_t))^{-1} g \rangle$$

Following from the definition of s_t in 1, we have $\|\psi'(\lambda_t)\|_{\Psi_t^*}^2 \leq \langle g_t, \text{diag}(s_t)^{-1} g_t \rangle$. Thus we have the following implication

$$\sum_{t=1}^T \|\psi'(\lambda_t)\|_{\Psi_t^*}^2 \leq \sum_{i=1}^d \|g_{1:T,i}\|_2$$

which prove the regret term (2).

Consequently, it remains to prove the bound over the Bregman divergence and the term $\Psi_T(\lambda^*)$ of proposition 2.1 and 2.2. Focusing on the composite mirror descent, we have:

$$\begin{aligned} B_{\Psi_{t+1}}(\lambda^*, \lambda_{t+1}) - B_{\Psi_t}(\lambda^*, \lambda_{t+1}) &= \frac{1}{2} \langle \lambda^* - \lambda_{t+1}, \text{diag}(s_{t+1} - s_t)(\lambda^* - \lambda_{t+1}) \rangle \\ &\leq \frac{1}{2} \max_i (\lambda_i^* - \lambda_{t+1,i})^2 \|s_{t+1} - s_t\|_1 \end{aligned}$$

Since $\|s_{t+1} - s_t\|_1 = \langle s_{t+1} - s_t, 1 \rangle$ and $\langle s_T, 1 \rangle = \sum_{i=1}^d \|g_{1:T,i}\|_2$ we have

$$\begin{aligned} \sum_{t=1}^{T-1} B_{\Psi_{t+1}}(\lambda^*, \lambda_{t+1}) - B_{\Psi_t}(\lambda^*, \lambda_{t+1}) &\leq \frac{1}{2} \sum_{t=1}^{T-1} \|\lambda^* - \lambda_{t+1}\|_\infty^2 \langle s_{t+1} - s_t, 1 \rangle \\ &\leq \frac{1}{2} \max_{t \leq T} \|\lambda^* - \lambda_t\|_\infty^2 \sum_{i=1}^d \|g_{1:T,i}\|_2 - \frac{1}{2} \|\lambda^* - \lambda_1\|_\infty^2 \langle s_1, 1 \rangle \end{aligned}$$

which prove us the term (3).

Finally, we also have that

$$\Psi_T(\lambda^*) = \delta \|\lambda^*\|_2^2 + \langle \lambda^*, \text{diag}(s_T) \lambda^* \rangle \leq \delta \|\lambda^*\|_2^2 + \|\lambda^*\|_\infty^2 \sum_{i=1}^d \|g_{1:T,i}\|_2$$

which give us the bound (1).

Performing a few algebraic simplification lead us to a corollary which give us a more intuitive form of the regret bound. Assume that \mathcal{X} is a compact set and set $D_\infty = \sup_{\lambda \in \mathcal{X}} \|\lambda - \lambda^*\|_\infty$, and define

$$\gamma_T \triangleq \sum_{i=1}^d \|g_{1:T,i}\|_2 = \inf_s \left\{ \sum_{t=1}^T \langle g_t, \text{diag}(s)^{-1} g_t \rangle : \langle 1, s \rangle \leq \sum_{i=1}^d \|g_{1:T,i}\|_2, s \succeq 0 \right\}$$

Corollary 2.1. *Assume that D_∞ and γ_T are defined as above. For $\{\lambda_t\}$ generated using the (3) with $\eta = \|\lambda^*\|_\infty$, for any $\lambda^* \in \mathcal{X}$ we have*

$$R_\phi(T) \leq 2\|\lambda^*\|_\infty \gamma_T + \delta \frac{\|\lambda^*\|_2^2}{\|\lambda^*\|_\infty} \leq 2\|\lambda^*\|_\infty \gamma_T + \delta \|\lambda^*\|_1$$

Using update (4) to generate $\{\lambda_t\}$ and setting $\eta = D_\infty/\sqrt{2}$, we have

$$R_\phi(T) \leq \sqrt{2}D_\infty \sum_{i=1}^d \|g_{1:T,i}\|_2 = \sqrt{2}D_\infty \gamma_T$$

All these presented results are respected by our problem. Indeed, as we stated before, our original problem is for sure a convex problem (a quadratic function with positive semidefinite Q) and also the constraints are convex (a set of disjoint unitary simplices) implying strong duality. As we know from theory, the dual problem is for sure convex, even if the original one it's not. What we should argue is the presence of the constraints on the dual variables. Being ADAGRAD applicable to any convex set $\mathcal{X} \subseteq \mathbb{R}^n$, this is suitable for our problem, since $\mathcal{X} = \{\lambda \geq 0\}$. We will use one of the just explained update, (preferably the *primal-dual* one, since we have a general proof of convergence over exactly \mathbb{R}_+^n , in LEMMA 2 of Appendix A of [4]), compute the update and then project it on the λ 's constraint set. We expect to achieve an asymptotically sub-linear regret, as the authors showed in their work.

2.5 General application to our case

In our setting, we aim to solve problem (P) using (D). We will use $\mathcal{X} = \{\lambda \geq 0\}$ as the constraint set of the lagrangian multipliers λ 's, and $\mathcal{Y} = \{\sum_{i \in I^k} x_i = 1 \ \forall k \in K\}$ as the constraint set of the primal variables.

To solve the problem (3) and (4) in our problem settings, we consider the update of the dual variable λ . Referring to the problem (D), we can freely choose among three different update rules, in order (1), (3) and (4):

$$\begin{aligned} \lambda_{t+1} &= P_{\mathcal{X}}\{\lambda_t + \eta \text{diag}(G_t)^{-1/2} g_t\} \\ \lambda_{t+1} &= \arg \max_{\lambda \in \mathcal{X}} \{\eta \langle \bar{g}_t, \lambda \rangle + \frac{1}{t} \Psi_t(\lambda)\} \\ \lambda_{t+1} &= \arg \max_{\lambda \in \mathcal{X}} \{\eta \langle g_t, \lambda \rangle + B_{\Psi_t}(\lambda, \lambda_t)\} \end{aligned} \tag{9}$$

We need to find a maximum for the considered update function and then project it onto the constraint set \mathcal{X} .

About the terms Ψ and B_Ψ , these will be replaced, according to what we have said before, respectively by:

$$\begin{aligned} \Psi_t(\lambda) &= \frac{1}{2} \langle \lambda, H_t \lambda \rangle \\ B_{\Psi_t}(\lambda, \lambda_t) &= \Psi_t(\lambda) - \Psi_t(\lambda_t) - \langle \nabla \Psi_t(\lambda_t), \lambda - \lambda_t \rangle \end{aligned}$$

Regarding the maximum, this can be simply solved by differentiating the function with respect to the variable to be maximized, and setting it equal to 0. The detailed derivation of each update can be found in the appendix A.

$$\hat{\lambda}_{t+1} = \lambda_t + \eta \text{diag}(G_t)^{-1/2} g_t \tag{Update (5)}$$

$$\hat{\lambda}_{t+1} = -H_t^{-1} t \eta \bar{g}_t \tag{Update (3)}$$

$$\hat{\lambda}_{t+1} = \lambda_t + \eta H_t^{-1} g_t \tag{Update (4)}$$

For the stepsize η , this can be fixed apriori and we will see being it crucial for the convergence of the algorithm. Different stepsize rules can be employed [7]

$\eta = h$	Constant step size rule, with $h > 0$
$\eta = \frac{h}{\ g^t\ }$	Constant step length
$\eta = \frac{\alpha}{\beta + t}$	Square summable but not summable, with $\alpha > 0$ and $\beta \geq 0$
$\eta = \frac{\alpha}{\sqrt{t}}$	Nonsummable diminishing
$\eta = \frac{f(x^*) - \phi(\lambda_t)}{\ g^t\ ^2}$	Polyak stepsize

For the dual variables, after we obtain $\hat{\lambda}_{t+1}$, we use the projection over the nonnegative orthant to get λ_{t+1} , formally

$$\lambda_{t+1} = P_{\mathcal{X}}(\hat{\lambda}_{t+1}) = \max\{0, \hat{\lambda}_{t+1}\}$$

which is a trivial problem tackled many times in the literature [8].

We need also to derive an update for the primal variables, which are needed at each iteration in order to compute the dual function $\psi(\lambda)$ and the subgradient $\nabla_{\lambda}\psi(\lambda)$. Given the previous multiplier value λ_t

$$x_{t+1} = \arg \min_{x \in \mathcal{Y}} \{x^T Q x + q^T x + \langle \lambda_t, x \rangle\}$$

We need to solve the lagrangian relaxation of a convex quadratic problem, depending on equality constraints made up of disjoint simplices. The constraint set \mathcal{Y} can be rewritten in a matrix form $Ax = b$, where the vector b is a $k \times 1$ vector of all ones, and the matrix A can be derived with the following simple algorithm:

Algorithm 2 Construct matrix A

```

procedure CONSTRUCT_A( $K, [I^k]$ )
   $A \leftarrow []$  ▷ Initialize  $k \times n$  empty matrix
  for  $k \leftarrow 1$  to  $K$  do
     $a_k \leftarrow \text{zeros}(n, 1)$  ▷  $n \times 1$  empty row vector
    for  $i \leftarrow 1$  to  $n$  do
      if  $i \in I^k$  then ▷ Check if index  $i$  is in the set  $I^k$ 
         $a_k[i] \leftarrow 1$ 
      end if
    end for
     $A[k, :] \leftarrow a_k$ 
  end for
end procedure

```

Using the **KKT** (Karush-Kuhn Tucker) conditions, we can solve the problem directly through linear algebra by constructing the following linear system:

$$\begin{cases} Qx + q - \lambda_t + \mu A = 0 \\ Ax - b = 0 \end{cases} \implies \begin{cases} Qx + \mu A = \lambda_t - q \\ Ax = b \end{cases} \implies \begin{bmatrix} Q & A^T \\ A & 0 \end{bmatrix} \begin{bmatrix} x \\ \mu \end{bmatrix} = \begin{bmatrix} \lambda_t - q \\ b \end{bmatrix}$$

a linear system that can be solved directly and efficiently by computing the pivoted LU factorization of the matrix and save it for the next iterations. Hence, using the LU factor-

ization:

$$\begin{aligned} \begin{bmatrix} Q & A^T \\ A & 0 \end{bmatrix} \begin{bmatrix} x \\ \mu \end{bmatrix} &= \begin{bmatrix} \lambda_t - q \\ b \end{bmatrix} \implies \begin{bmatrix} x \\ \mu \end{bmatrix} = \begin{bmatrix} Q & A^T \\ A & 0 \end{bmatrix}^{-1} \begin{bmatrix} \lambda_t - q \\ b \end{bmatrix} \implies \\ &\begin{bmatrix} x \\ \mu \end{bmatrix} = U^{-1} * \left(L^{-1} * \left(P * \begin{bmatrix} \lambda_t - q \\ b \end{bmatrix} \right) \right) \end{aligned}$$

which can be solved efficiently through forward substitution (for the matrix L) and back substitution (for the matrix U) with a complexity of $O((n + K)^2)$.

Lastly, regarding the stopping condition for the algorithm termination, we can check the dual variables residual. Follows that, given a certain value of tolerance ε , we have

$$\|\lambda_t - \lambda_{t-1}\|_2 \leq \varepsilon \quad (11)$$

In addition to this, what we can check too is the duality gap between the original function and the dual function (provided that we know the optimal value), so

$$\begin{aligned} f(x) - \psi(\lambda) \\ f(x) - f(x^*) \leq f(x) - \psi(\lambda) \end{aligned} \quad (12)$$

From this we understand that having a zero duality gap implies optimality of the solution, and so fixing a certain value of tolerance τ we have a second stopping condition

$$f(x^*) - \psi(\lambda_i) \leq \tau$$

for each λ_i computed during iterations.

The starting value of the λ 's iterates is a vector of ones and the starting x is computed immediately using the starting λ and the factorization explained before.

Putting everything together, we obtain a slightly different algorithm than 1:

Algorithm 3 ADAGRAD on our dual problem

```

function ADAGRAD( $\eta, \delta, \varepsilon, \tau, max\_iter$ )
   $\lambda_0 \leftarrow 1$ 
   $g_{1:0} \leftarrow []$ 
   $x_0 \leftarrow$  system solution
   $\phi = \phi(\lambda_0)$ 
  for  $t \leftarrow 1$  to  $max\_iter$  do
     $g_t \leftarrow \frac{\partial \psi_\lambda(\lambda)}{\partial \lambda}$  ▷ Compute subgradient B
     $g_{1:t} \leftarrow [g_{1:t-1} \ g_t]$  ▷ Store subgradient
     $s_{t,i} \leftarrow \|g_{1:t,i}\|_2$  ▷ Solution of the problem 2.4
     $H_t \leftarrow \delta I + diag(s_t)$ 
     $x_t = \arg \min_{x \in \mathcal{Y}} \{x^T Q x + q x + \langle \lambda_{t-1}, x \rangle\}$  ▷ Lagrangian
     $\hat{\lambda}_t =$  one among (9)
     $\lambda_t = P_{\mathcal{X}}(\hat{\lambda}_t)$ 
    if  $\phi(\lambda_t) \geq \phi$  then
       $\phi = \phi(\lambda_t)$ 
    end if
    if (11) OR (12) is true then
      return  $\phi$ 
    end if
  end for
end function

```

3 Implementation

The ideas above described have been implemented using Julia. The provided package give us the ability to provide the dimension of the problem n and the number of simplices K .

Subsequently, the code is all automated and provide the creation of the matrix $Q \succeq 0$, A , the random sets I^k , the vector q and all the required parameters. The update rules that you want to test can be chosen as well as the different stepsize rules that you want to employ. The structures created are stored in *mat* file and loaded each time to avoid time consumption in the creation of the different objects and also for the derivation of the optimal primal solution. Indeed this latter is found using the off-the-shelf solver **YALMIP** [9], available on **MATLAB**, where the generated *mat* file is stored and a simple script is used to derive the primal solution.

Also, there are some hyperparameters which can be modified by the user inside *main.jl*, like the maximum number of iterations and the ε value to check the λ residual.

All the values assumed by each λ -norm and x -norm residual, the number of iterations, the value of the dual function and the gap at the current iteration are stored in a *csv* file. The last one is used then to generate some plots showing the convergence towards the optimal solution.

The customized **ADAGRAD** display every ten iterations the time elapsed, the ϕ value, both the λ norm residual and the x one, and the dual gap $f(x^*) - \phi(\lambda)$.

In order to test the program on different problem dimensions, we have chosen different sizes to perform our experiments, in order to have a "grid search" of different experimentations:

n	K	ε	τ	α	h	δ	max_iter
1000	[20]	10^{-14}	10^{-7}	0.8	[]	10^{-16}	100000
5000	[10]		10^{-6}	0.3	[200, 25, 5, 2, 10^{-2} , 10^{-3}]		

The code was executed using Julia *REPL* tool. It has a slowish startup time, so the first execution are "warm-up" execution, but as soon as the entire code is compiled, all modules are loaded and the needed memory is computed. Q is the resulting covariance matrix of a random generated matrix A_h , in formula:

$$Q = A_h^T * A_h$$

which we know being symmetric and surely positive semidefinite.

Regarding the factorization, the entire KKT matrix is factorized using the *Julia* method `lu!()`.

Another observation regard the subgradient: being the ϕ non differentiable in the general case, the direction is not guaranteed to be ascent/descent. So we keep track of the best $\phi(\lambda)$, best iteration, best x and λ to get the suboptimal solution found. This sometimes result in the dual function being minimized instead of maximized. In certain cases we observed also that when we reach a negative dual gap (due to big steps and not feasibility of the actual iterates) changing the direction of the subgradient brings improvement in reaching the solution.

Last but not least, we noticed in the experiments that the stepsize selection is crucial for the convergence of the algorithm. In particular, the value of the hyperparameters α , β and δ should be chosen ad-hoc to guarantee that we have a "smooth" optimization.

What we noticed in particular is about the different update rules and hyperparameters. As already said, the value of α , β and δ are fundamental, the latter for the update rules (3) and (4) while the first two for the standard (5).


In the first two mentioned updates, the value of H_t^{-1} acts as a sort of "regularizer", scaling in both cases the value of the gradient. Modifying δ impacts primarily on the steps taken, with smaller value of it resulting in bigger steps while a low value of it result in a very low convergence (even too much).

Referring always to these two updates, the only stepsize rule that give a "fast" convergence with big steps is the first one, a constant stepsize with a variable value of h . We change the value of h during iteration making it smaller as soon as we reach the optimal gap.

The rule (3) shows no improvement even adjusting these parameters, while rule (4) converges slowly with respect to the standard one.

Indeed it is rule (5) that give us the best results: we are capable of reaching the desired dual gap in fewer iterations, adjusting the value of the hyperparameters. In this case, we used the nonsummable diminishing rule fixing the value of α , switching to Polyak stepsize in the tail.

We can show some plots and timing results of the experiments reported in the table:



<i>Total iterations</i>	3092
<i>total time (sec)</i>	665.3 s
<i>Best $\phi(\lambda) - f(x^*)$</i>	1.414e-06
<i>Best iteration</i>	3092
<i>Best $\psi(\lambda)$</i>	1.202e+05
<i>Best $\ \lambda_t - \lambda_{t-1}\$</i>	2.586e-08
<i>Best $\ x_t - x_{t-1}\$</i>	5.190e-11

Table 1: Sum up table

Figure 1: Dual gap with update rule 1, $n = 5000$ and $K = 10$

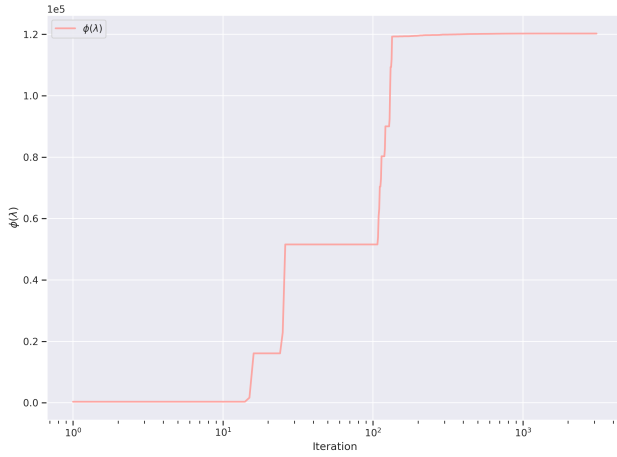


Figure 2: Dual value with update rule 1, $n = 5000$ and $K = 10$


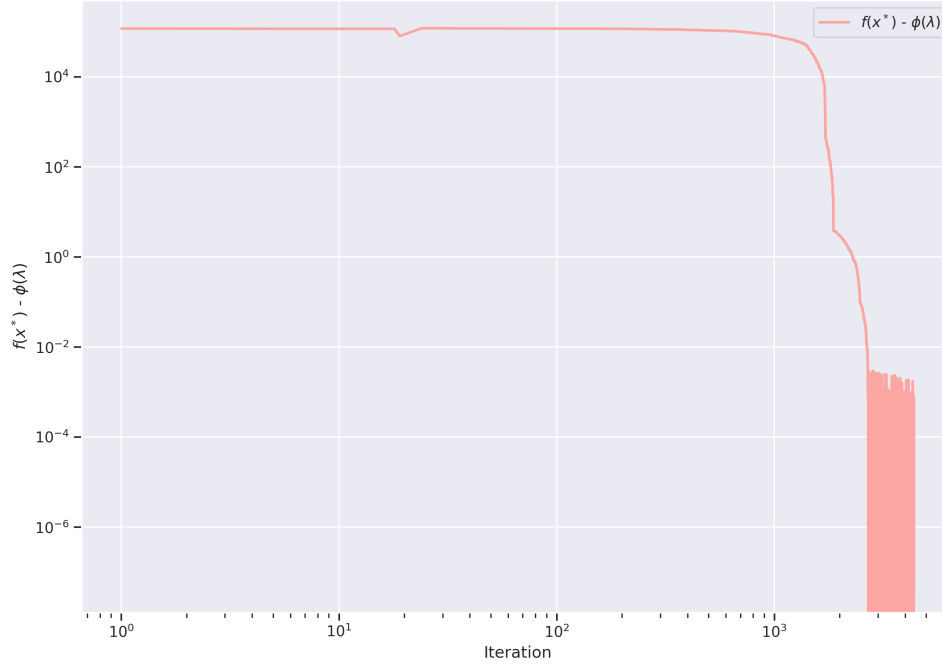


Figure 3: Norms with update rule 1, with $n = 5000$ and $K = 10$

As we can see by the results of the update rule 1, the slope of the curve change in the tail, reaching an optimal dual gap. This is due to the big function value we have and the stepsize we employ. In the tail, we can see how the value start oscillating near the optimal achieving the desired dual gap after a few iterations.

The same holds for the norms of the variables, oscillating but achieving an almost zero change in the tail.



<i>Total iterations</i>	4379
<i>total time (sec)</i>	952.5 s
<i>Best $\phi(\lambda) - f(x^*)$</i>	5.243e-08
<i>Best iteration</i>	4379
<i>Best $\psi(\lambda)$</i>	1.202e+05
<i>Best $\ \lambda_t - \lambda_{t-1}\$</i>	1.471e-09
<i>Best $\ x_t - x_{t-1}\$</i>	4.153e-05

Table 2: Sum up table

Figure 4: Dual gap with update rule 3, $n = 5000$ and $K = 10$

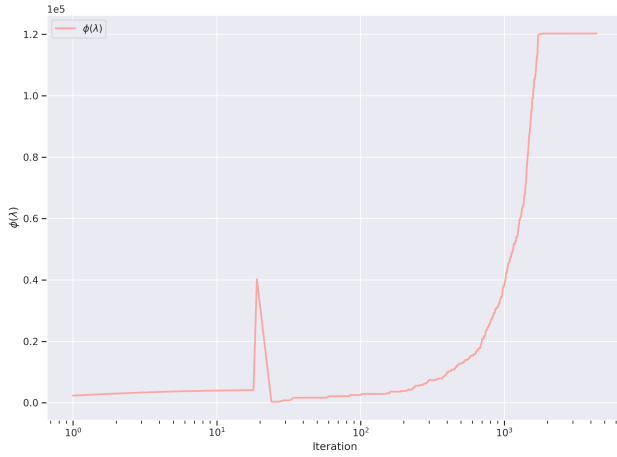


Figure 5: Dual value with update rule 3, $n = 5000$ and $K = 10$




Figure 6: Norms with update rule 3, with $n = 5000$ and $K = 10$

As we can see by the results of the update rule 3, here we perform different steps, by adjusting the value of h during the optimization. Indeed the slope of the curve change also in the end, this time with a more high value. The same happens also in this case, we start oscillating around the optimal value until we get the desired dual gap.

The norms too oscillate around very small values.

4 Code description

The code is provided in Julia through package. To try it, simply open the Julia REPL inside the current folder, and launch `julia` specifying the option

```
--project=.
```

then enter in the `pkg` mode using `]` and launch

```
instantiate
```

to download the required modules. Then go back to the command line tool (simply backspace) and run

```
include(src/main.jl)
```

At this point you will be asked for a value of n and K .

All the code is contained inside the `src` folder:

- *Utils.jl*: module containing some useful functions:
 - `construct_full_matrix(Q, A, K)`: return the entire KKT matrix;
 - `construct_A(K, n, I_K)`: return the constraint matrix A as described in the report;
- *ADAGRAD_Solver.jl*: module containing:
 - a struct `Solver`, containing all the parameters and return value needed for the problem;
 - functions to compute the λ/x -norm, the subgradient, the γ value and the update rules;
 - `my_ADAGRAD`, which implement the algorithm derived in the report;
- *main.jl*: used to test all the code;
- *solve_prob.m*: MATLAB script to compute optimal primal solution using YALMIP;

In addition, the entire project folder contains:

- *papers*: folder containing all the literature referenced in the bibliography;
- *results*: folder containing all the plots and logs obtained with the experimentations values of table 3;
- *util.sh*: a clean-up routine for the *logs* and *plots* produced;
- *analyzer.py*: used to inspect the collected *.csv* files and get some stats;
- *plotter.py*: used to generate plots from the *.csv* files;

References

- [1] John Duchi, Elad Hazan, and Yoram Singer. Adaptive subgradient methods for online learning and stochastic optimization. *Journal of Machine Learning Research*, 12(61):2121–2159, 2011.
- [2] Lin Xiao. Dual averaging method for regularized stochastic learning and online optimization. In Y. Bengio, D. Schuurmans, J. Lafferty, C. Williams, and A. Culotta, editors, *Advances in Neural Information Processing Systems*, volume 22. Curran Associates, Inc., 2009.
- [3] John Duchi, Shai Shalev-Shwartz, Yoram Singer, and Ambuj Tewari. Composite objective mirror descent. In *Composite Objective Mirror Descent*, pages 14–26, 12 2010.
- [4] Antonio Frangioni, Bernard Gendron, and Enrico Gorgone. On the computational efficiency of subgradient methods: a case study with lagrangian bounds. *Mathematical Programming Computation*, 9(4):573–604, Dec 2017.
- [5] Neal Parikh and Stephen Boyd. Proximal algorithms. *Foundations and Trends® in Optimization*, 1(3):127–239, 2014.
- [6] Yurii Nesterov. Primal-dual subgradient methods for convex problems. *Mathematical Programming*, 120(1):221–259, Aug 2009.
- [7] Marina A. Epelman. Ioe 511/math 652: Continuous optimization methods, section 1, 2007. Lecture notes =<http://www-personal.umich.edu/~mepelman/teaching/511notesFA07.pdf>.
- [8] Ang Andersen. Projection onto nonnegative orthant, rectangular box and polyhedron, 2020. First draft: March 19, 2020; Last update: December 23, 2020. Université de Mons, https://angms.science/doc/CVX/Proj_nonnegBoxpoly.pdf.
- [9] J. Löfberg. Yalmip : A toolbox for modeling and optimization in matlab. In *In Proceedings of the CACSD Conference*, Taipei, Taiwan, 2004.

A Appendix A: Updates derivations

Given a matrix A of size $n \times n$, the proximal term expression of a vector x of size $n \times 1$ at step t is

$$\Psi_t(x) = \frac{1}{2} \langle x, Ax \rangle$$

A.1 Differentiating proximal term

Derivation of proximal term arises from the need of obtaining the $\arg \max x$ of a given update function. Starting by the definition

$$\frac{\partial \Psi_t(x)}{\partial x} = \frac{1}{2} \frac{\partial}{\partial x} \langle x, Ax \rangle$$

which can be rewritten in a more classical form, and the derivative is immediate

$$\frac{1}{2} \frac{\partial}{\partial x} x^T Ax = Ax$$

A.2 Derivation of primal-dual update

We can now look into the detailed derivation of the update (3). Starting from the equation

$$\lambda_{t+1} = \arg \max_{\lambda \in \mathcal{X}} \{ \eta \langle \bar{g}_t, \lambda \rangle + \frac{1}{t} \Psi_t(\lambda) \}$$

we want to achieve the $\arg \max$ over the set \mathcal{X} . Focusing on the maximum problem, we want to get the maximum λ , hence:

$$\begin{aligned} \frac{\partial}{\partial \lambda} \eta \langle \bar{g}_t, \lambda \rangle + \frac{\partial}{\partial \lambda} \frac{1}{t} \Psi_t(\lambda) &= 0 \\ \eta \bar{g}_t + \frac{1}{2t} \frac{\partial}{\partial \lambda} \langle \lambda, H_t \lambda \rangle &= 0 \end{aligned}$$

Knowing that H_t is a diagonal matrix and using the previous derivation:

$$\frac{\partial}{\partial \lambda} \langle \lambda, H_t \lambda \rangle = 2 H_t \lambda$$

Substituting this into the above derivation, we get the maximum λ :

$$\begin{aligned} \eta \bar{g}_t + \frac{1}{t} H_t \lambda &= 0 \\ \lambda &= -H_t^{-1} t \eta \bar{g}_t \end{aligned}$$

A.3 Derivation of composite-mirror update

We follow the same approach also for the update (4). Starting from the definition

$$\lambda_{t+1} = \arg \max_{\lambda \in \mathcal{X}} \{ \eta \langle g_t, \lambda \rangle + B_{\Psi_t}(\lambda, \lambda_t) \}$$

where we remark the definition of Bregman divergence

$$B_{\Psi_t}(\lambda, \lambda_t) = \Psi_t(\lambda) - \Psi_t(\lambda_t) - \langle \nabla \Psi_t(\lambda_t), \lambda - \lambda_t \rangle$$

Also in this case, to find the $\arg \max$ we derive the update formula and set it to be zero

$$\begin{aligned} \frac{\partial}{\partial \lambda} \eta \langle g_t, \lambda \rangle + \frac{\partial}{\partial \lambda} B_{\Psi_t}(\lambda, \lambda_t) &= 0 \\ \eta g_t + \frac{\partial}{\partial \lambda} [\Psi_t(\lambda) - \Psi_t(\lambda_t) - \langle \nabla \Psi_t(\lambda_t), \lambda - \lambda_t \rangle] &= 0 \\ \eta g_t + H_t \lambda - \frac{\partial}{\partial \lambda} \langle \nabla \Psi_t(\lambda_t), \lambda - \lambda_t \rangle &= 0 \end{aligned}$$

Considering the first term of the scalar product, we have to evaluate the first order taylor model of the proximal function Ψ_t at the known quantity λ_t . It is easy to see that by using the derivation in the first paragraph and the fact that H_t is diagonal

$$\nabla \Psi_t(\lambda_t) = \nabla \frac{1}{2} \langle \lambda_t, H_t \lambda_t \rangle = H_t \lambda_t$$

Putting this into equation

$$\frac{\partial}{\partial \lambda} \langle H_t \lambda_t, \lambda - \lambda_t \rangle = \frac{\partial}{\partial \lambda} \sum_{i=1}^n h_{ii} \lambda_{t,i} (\lambda_i - \lambda_{t,i})$$

And so is easy to see that

$$\frac{\partial}{\partial \lambda} \langle H_t \lambda_t, \lambda - \lambda_t \rangle = \begin{bmatrix} \frac{\partial}{\partial \lambda_1} \langle H_t \lambda_t, \lambda - \lambda_t \rangle \\ \frac{\partial}{\partial \lambda_2} \langle H_t \lambda_t, \lambda - \lambda_t \rangle \\ \dots \\ \frac{\partial}{\partial \lambda_n} \langle H_t \lambda_t, \lambda - \lambda_t \rangle \end{bmatrix} = \begin{bmatrix} h_{11} \lambda_{t,1} \\ h_{22} \lambda_{t,2} \\ \dots \\ h_{nn} \lambda_{t,n} \end{bmatrix} = H_t \lambda_t$$

And finally putting this into the main problem, we get:

$$\begin{aligned} \eta g_t + H_t \lambda - H_t \lambda_t &= 0 \implies H_t \lambda = H_t \lambda_t - \eta g_t \\ \lambda &= H_t^{-1} [H_t \lambda_t - \eta g_t] = \lambda_t + \eta H_t^{-1} g_t \end{aligned}$$

B Appendix B: subgradient computation

Utterly following the definition, we state that s is a subgradient of ψ at x

$$\psi(\lambda) \leq \psi(x) + s^T (\lambda - x) \quad \forall \lambda \in \mathbb{R}^n$$

which reordering the term can be written as

$$s^T (\lambda - x) \geq \psi(\lambda) - \psi(x) \quad \forall \lambda \in \mathbb{R}^n$$

Either we can solve the above complex inequality with unknown variables s and λ , or better we can apply the following reasoning. The subdifferential is the set containing all the numbers in the interval $[a, b]$ such that

$$\begin{aligned} a &= \lim_{\lambda \rightarrow \lambda_0^-} \frac{\psi(\lambda) - \psi(\lambda_{t-1})}{\lambda - \lambda_{t-1}} \\ b &= \lim_{\lambda \rightarrow \lambda_0^+} \frac{\psi(\lambda) - \psi(\lambda_{t-1})}{\lambda - \lambda_{t-1}} \end{aligned}$$

where the $^+$ and $^-$ portion must have a certain value greater than zero (order of 10^{-2}).